

⑨

MDP: Part 2

* Convergence

Case 1: If the tree has maximum depth M , then V_M holds the actual untruncated values.

Case 2: If the discount is less than 1


* Policy Evaluation

⇒ Given the policy, you want to know, how good is that policy.

* Utilities for a fixed Policy

$V^\pi(s)$ = Expected total discounted rewards starting in s and following π

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$



⇒ How do we calculate the V 's for a fixed policy π ?

→ Idea 1: Turn recursive Bellman equations into update.

→ Idea 2: Without the maxes, the Bellman equations are just a linear system.

* Computing Actions from Values

$$\pi^*(s) = \underset{\text{argmax } a}{\left(\sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \right)}$$

↳ Policy Extraction

* Computing Actions from Q-Values

$$\pi^*(s) = \underset{a}{\text{argmax}} Q^*(s, a)$$

"Actions are easier to select from"
q-values than values!

* Problems with Value Iteration

Problem 1: It's slow - $O(S^2A)$ per iteration

Problem 2: The "max" at each state rarely changes

Problem 3: The policy often converges long before the value.

* Policy Iteration

Step 1: (Policy Evaluation)

Calculate utilities for some fixed policy (not optimal utilities)
until convergence

Step 2: (Policy Improvement)

Update policy using one-step look-ahead with resulting converged (but not optimal) utilities as future values

⇒ Repeat Steps until policy Converges.

$$\textcircled{1} \quad V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

$$\textcircled{2} \quad \pi_{i+1}(s) = \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

⇒ Both Value Iteration and Policy Iteration are dynamic programs for solving MDPs

* Reinforcement Learning

⇒ Important ideas in reinforcement Learning:

* Exploration: You have to try unknown actions to get information

* Exploitation: Eventually, you have to use what you know

* Regret: Even if you learn intelligently, you make mistakes.

* Sampling: Because of chance, you have to try things repeatedly.

* Difficulty: Learning can be harder than solving a known MDP.