

## Reinforcement Learning II

### \* How to Explore?

⇒ Several Schemes for forcing exploration

# Simplest: random action ( $\epsilon$ -greedy)

- Every time step, flip a coin
- With (small) probability  $\epsilon$ , act randomly
- With (large) probability  $1-\epsilon$ , act on current policy

⇒ Problem with random actions?

- You do eventually explore the space, but keep thrashing around once learning is done.
- One Solution: lower  $\epsilon$  over time
- Another Solution: Exploration functions

### \* Exploration Function

- Random actions: Explore a fixed amount
- Better idea: Explore areas whose goodness is not (yet) established, eventually stop exploring

⇒ Takes a value estimate  $u$  and a visit count  $n$  & return an optimistic utility:

$$f(u, n) = u + \frac{k}{n}$$

Modified:  
Q-update:  $Q(s, a) \leftarrow_\alpha R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$



## ★ Regret

⇒ Measure of your total mistake cost

Difference between your rewards  
\* optimal rewards

$$\text{Regret} = U(\text{Optimal action}) - U(\text{action taken})$$



## \* Approximate Q-Learning

- ⇒ Basic Q-Learning keeps a table of all  $q$ -values.
- ⇒ In realistic situations we cannot possibly learn about every single state.
  - Too many states to visit them all in training.
  - Too many states to hold the  $q$ -tables in memory.
- Instead, we want to generalize:
  - Learn about some small number of training states from experience.
  - Generalize that experience to new, similar situations.
  - This is a fundamental idea in machine learning, and we will see it over & over again.

## # Feature-Based Representations

"Describe state using a vector of features"

- ⇒ Features are functions from states to real numbers (often 0/1) that capture important properties of the state.

→ Example:

- Distance to closest ghost
- Distance to closest dot
- Number of ghosts
- $1 / (\text{dist to dot})^2$
- Is pacman in a tunnel? (0/1)



## # Linear Value Function

⇒ Using feature representation, we can write a function (or value function) for any state using a few weights:

$$V(s) = \omega_1 f_1(s) + \omega_2 f_2(s) + \dots + \omega_m f_m(s)$$

$$Q(s) = \omega_1 f_1(s, a) + \omega_2 f_2(s, a) + \dots + \omega_m f_m(s, a)$$

Advantage: One experience is summed up in a few powerful numbers.

Disadvantage: States may share features but actually be very different in value.

⇒ Q-learning with linear Q functions:

$$\text{transition} = (s, a, r, s')$$

$$\text{difference} = [r + \gamma \max_{a'} Q(s', a)] - Q(s, a)$$

$$\text{Exact } Q: Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

$$\text{Approximate } Q: \omega_i \leftarrow \omega_i + \alpha [\text{difference}] f_i(s, a)$$

## # Q-Learning & Least Square

$$\text{let } y = r + \gamma \max_{a'} Q(s', a)$$

$$e_{\text{rmse}} = y - Q(s, a)$$

$$\rightarrow \omega_1 f_1 + \omega_2 f_2 + \dots + \omega_m f_m$$

$$\nabla_{\omega} \frac{1}{2} e^2 = \nabla_{\omega} (y - \omega^T f)^2 = -2(y - \omega^T f) f$$

$$\omega \leftarrow \omega + \alpha [\text{diff.}] f(s, a)$$



## \* Policy Search

⇒ Learn policies that maximize rewards, not the value that predict them.

⇒ Simplest policy search:

- ↳ Start with an initial value function or Q function
- ↳ Nudge each feature weight up & down & see if your policy is better than before

