

⑤

Adversarial Search

★ Types of Games

⇒ Axes

- Deterministic or Stochastic?
- One, two or more players?
- Zero Sum?
- Perfect information?
(Can you see the state)

⇒ Want algorithms for calculating a **Strategy (Policy)** which recommends a move from each state.

★ Deterministic Games

⇒ Many possible formalization, one is:

- States S (start at S_0)
- Players $P = \{1 \dots N\}$ (usually take turns)
- Actions: A (may depend on player / state)
- Transition Function: $S \times A \rightarrow S$
- Terminal Test: $S \rightarrow \{t, f\}$
- Terminal Utilities: $S \times P \rightarrow R$

⇒ Solution for a player is a **Policy**:
 $(S \rightarrow A)$

Zero-Sum Game

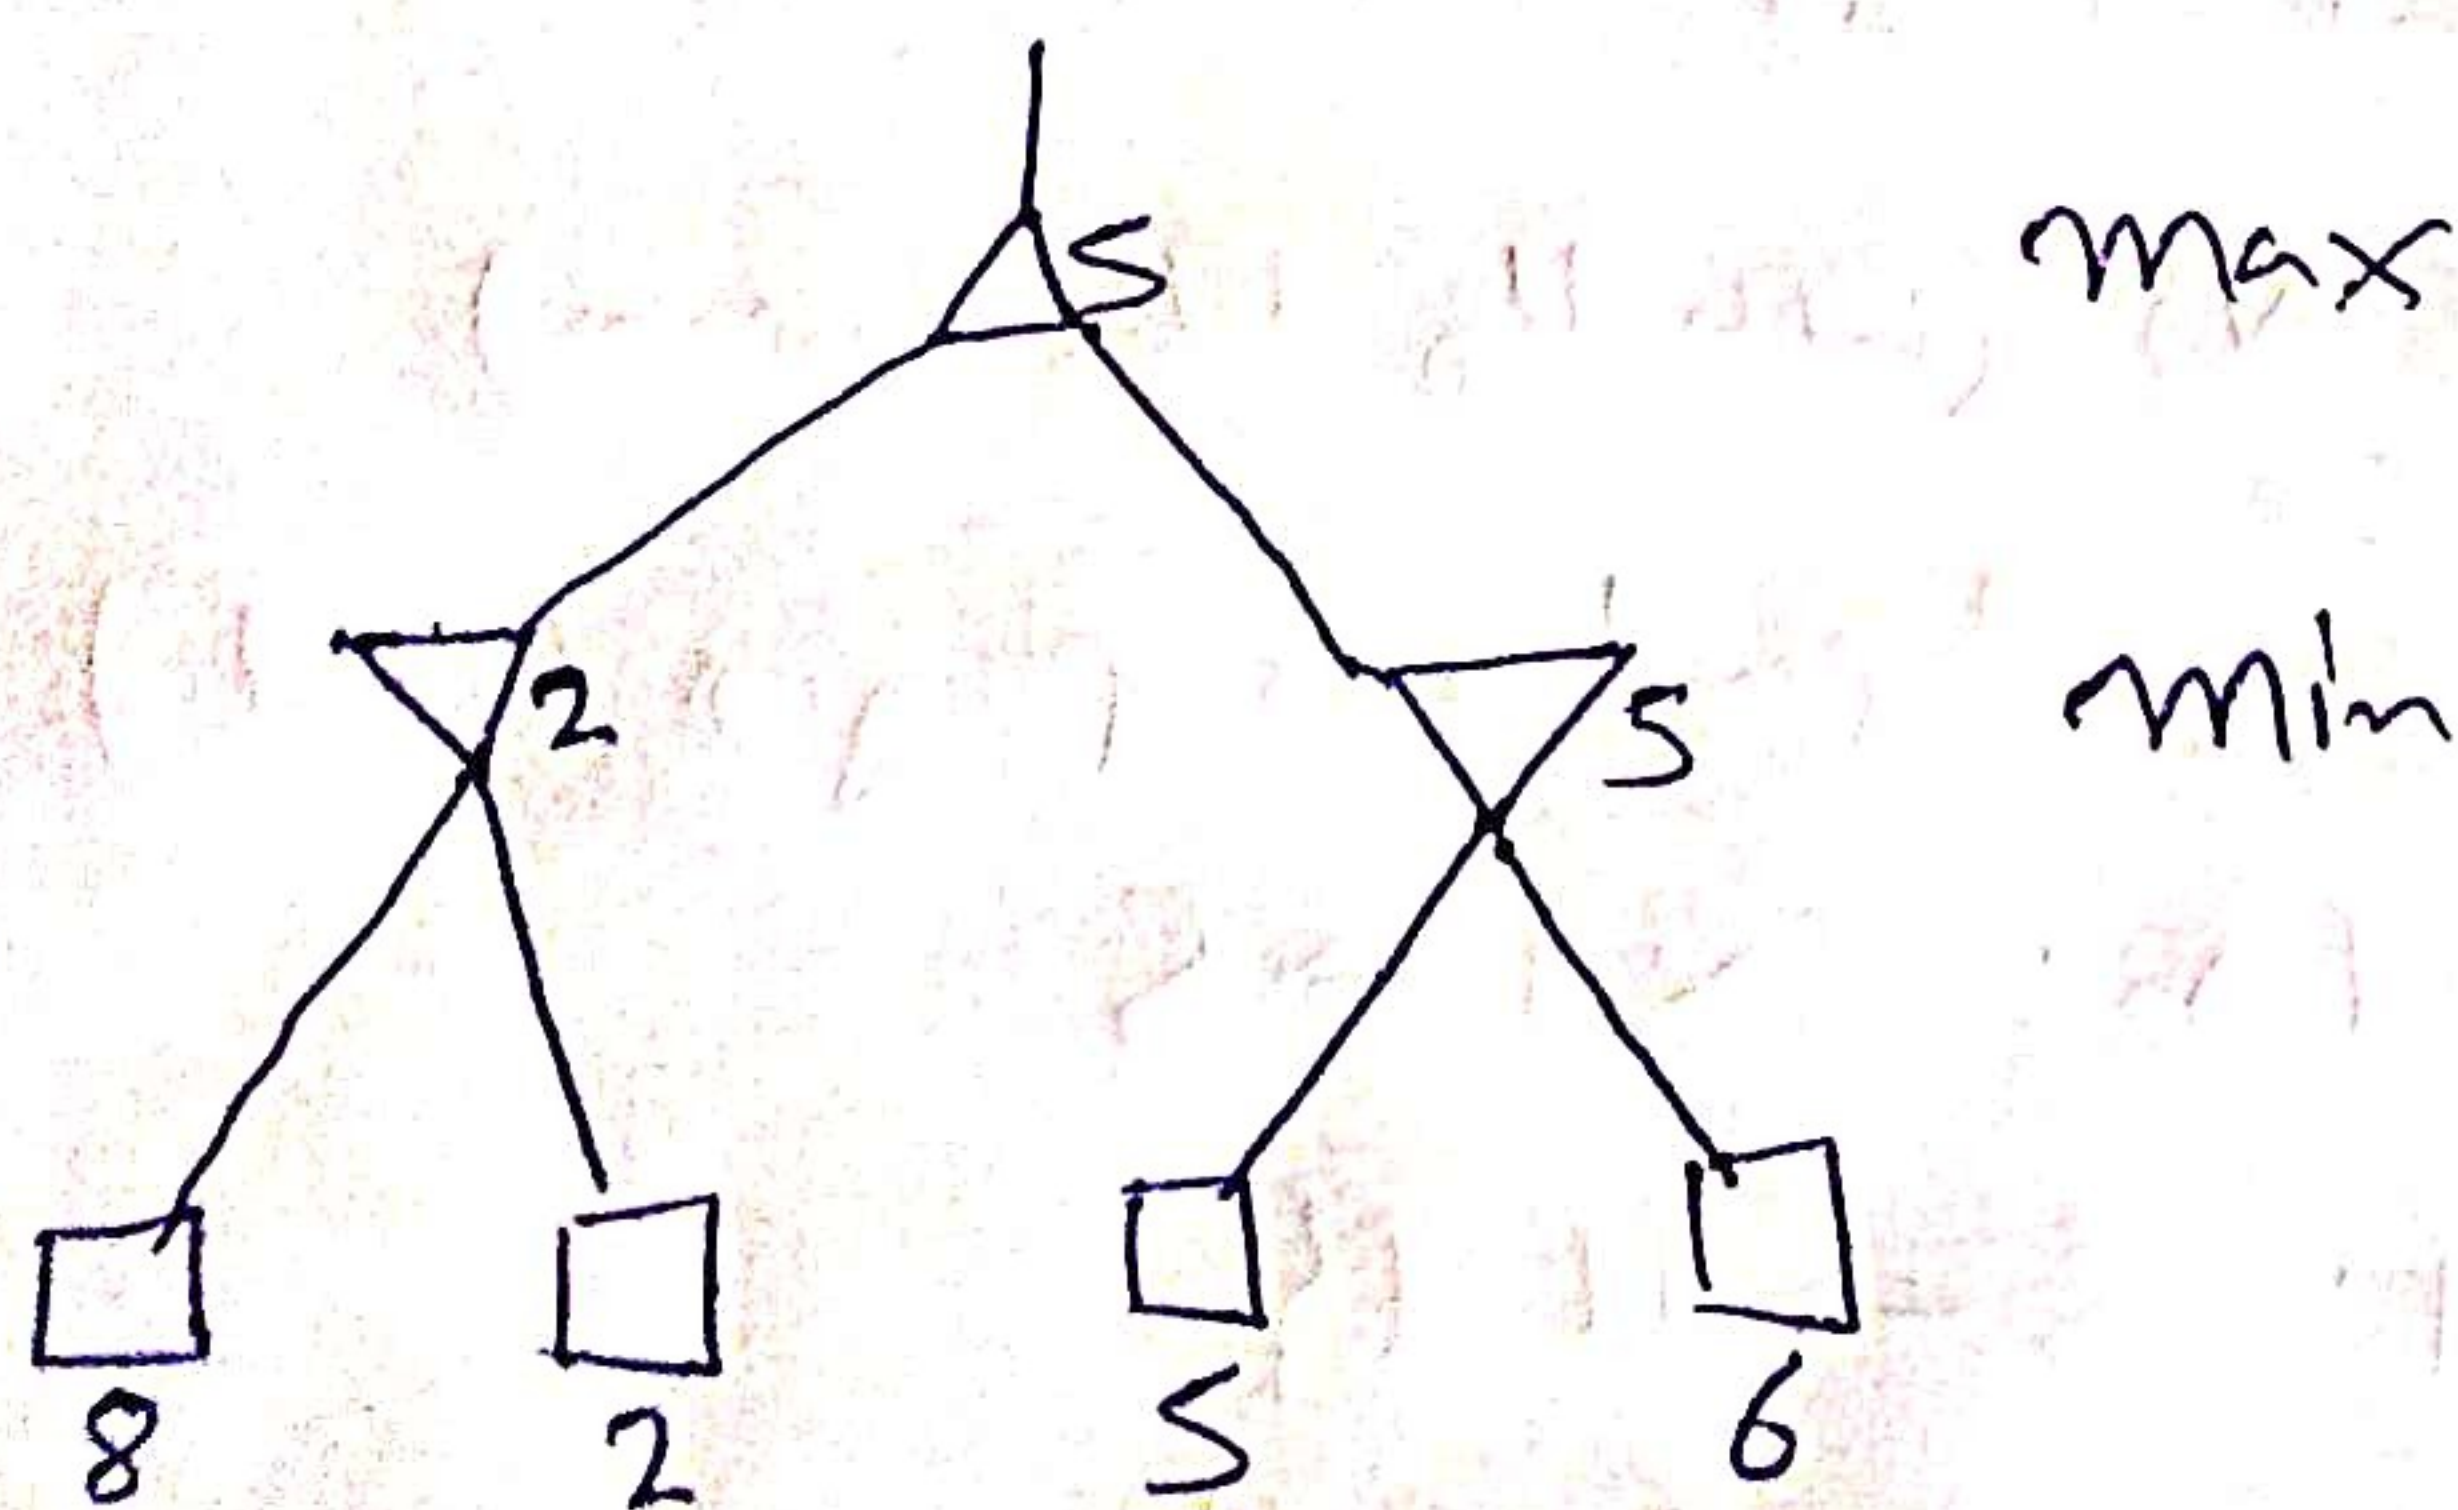
- ⇒ Agents have opposite utilities.
- ⇒ Single value that one maximizes and the other minimizes
- ⇒ Adversarial, pure competition

General Game

- ⇒ Agents have independent utilities

★ Adversarial Search (Minimax)

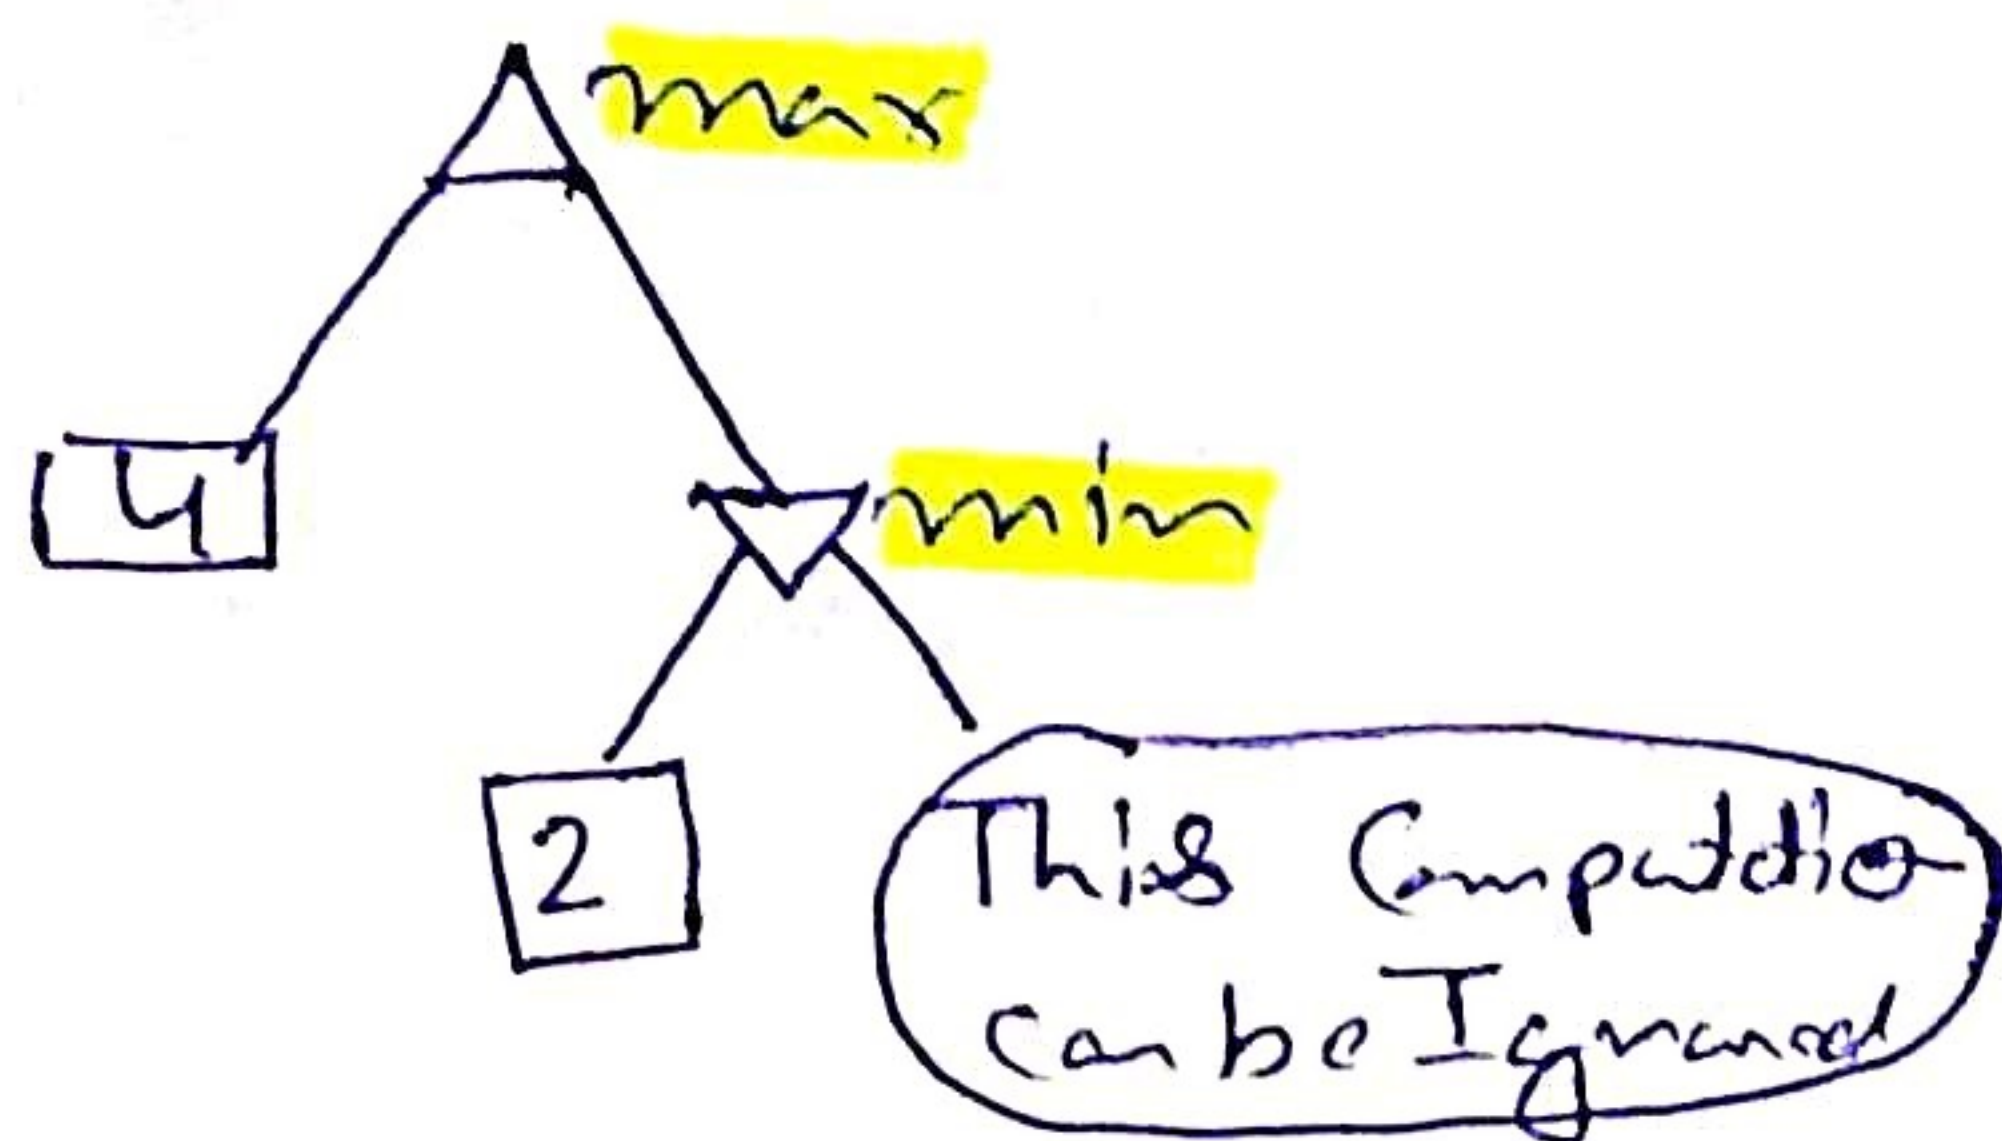
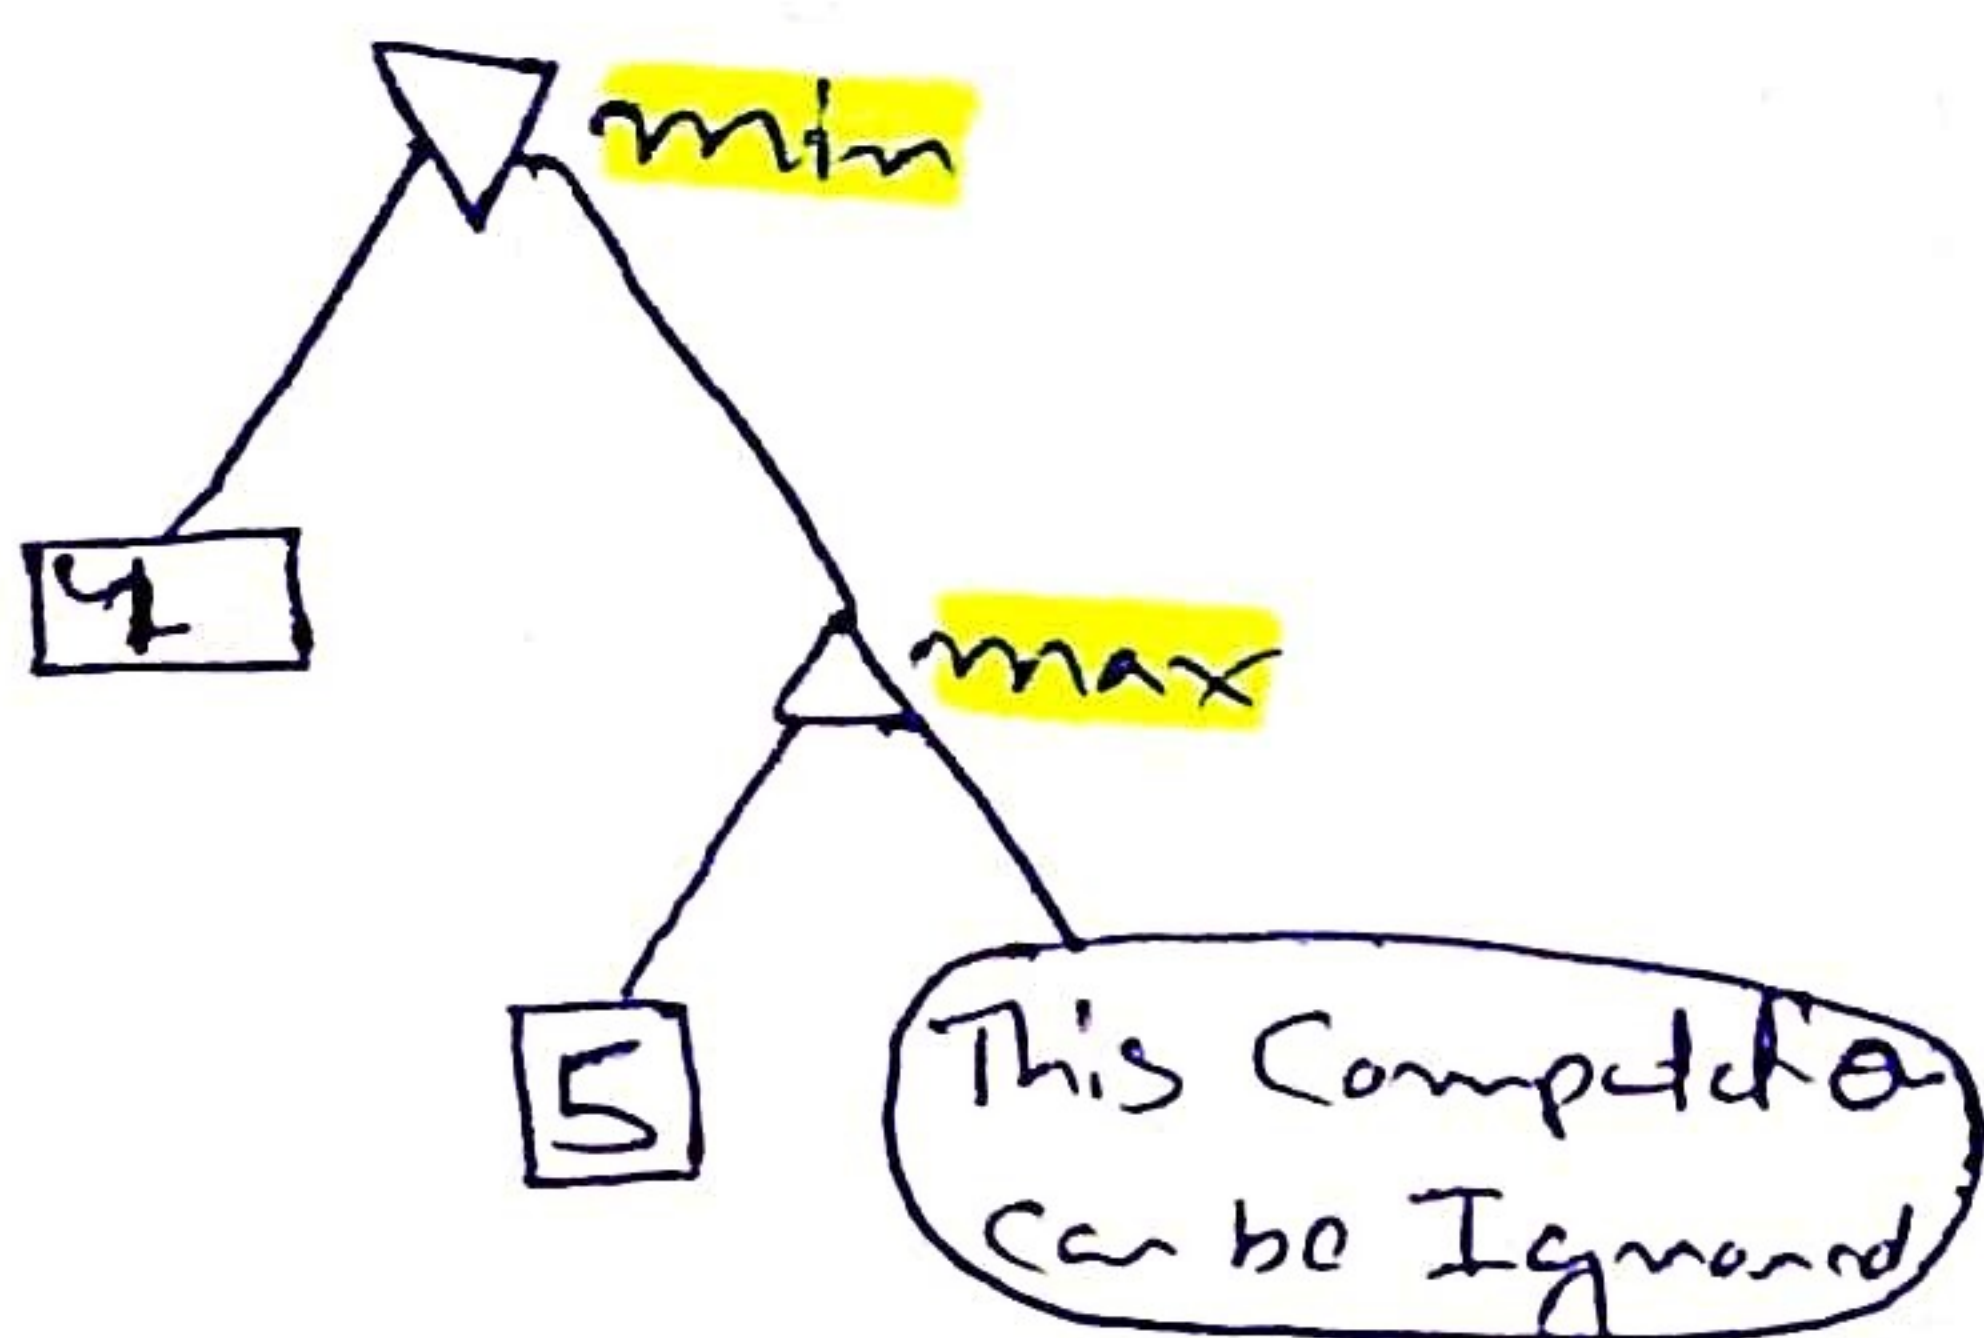
- A state-space search tree
- Players alternate turns
- Compute each node's minimax value:
{The best achievable utility against a}
{actional (optimal) adversary}



```
def max-value(state):  
    initialize  $V = -\infty$   
    for each successor of state:  
         $V = \max(V, \text{min-value}(\text{successor}))$   
    return  $V$ 
```

```
def min-value(state):  
    initialize  $V = +\infty$   
    for each successor of state:  
         $V = \min(V, \text{max-value}(\text{successor}))$   
    return  $V$ 
```


* Alpha-Beta Pruning



α : MAX's best option on path to root

β : MIN's best option on p

def max-value(State, α , β):
 initialize $V = -\infty$
 for each Successor of State:
 $V = \max(V, \text{Value}(\text{Successor}, \alpha, \beta))$
 if $V \geq \beta$ return V
 $\alpha = \max(\alpha, V)$
 return V

def min-value(State, α , β):
 initialize $V = +\infty$
 for each successor of State:
 $V = \min(V, \text{value}(\text{successor}, \alpha, \beta))$
 if $V \leq \alpha$ return V
 $\beta = \min(\beta, V)$
 return V

* Alpha-Beta pruning Properties

- ⇒ This pruning has no effect on minimax value computed for the root!
- ⇒ Values of intermediate nodes might be wrong.
- ⇒ Good child ordering improves effectiveness of pruning.

* Resource Limits

Problem: In realistic game, cannot search to leaves!

Solution: Depth-limited Search

- Search only to a limited depth in the tree.
- Replace terminal utilities with a evaluation function for non-terminal positions.

- ⇒ Guarantee of optimal play is gone
- ⇒ Use iterative deepening for an anytime algorithm.

* Depth Matters

- ⇒ Evaluation functions are always imperfect.
- ⇒ The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters.

* Synergies between Evaluation Function and Alpha-Beta

⇒ Alpha-Beta: Amount of pruning depends on expansion ordering.

↳ Evaluation function can provide guidance to expand most promising nodes first

