

Naive Bayes

→ Feature vectors x are discrete-valued.

Motivating Example: Spam filter
 ↗ Spam
 ↘ Not Spam

⇒ We'll begin our construction of our spam filter by specifying the features x_i : word to represent an email.

⇒ We will represent an email via a feature vector whose length is equal to the number of words in the dictionary.

↳ If an email contains the i^{th} word of the dictionary, then we will set $x_i = 1$; otherwise, we let $x_i = 0$.

⇒ The set of words encoded into the feature vector is called the **Vocabulary**, so the dimension of x is equal to the size of vocabulary.

⇒ If Vocabulary = 50000 words, $x \in \{0, 1\}^{50,000}$.

↳ If we were to model x explicitly with a multinomial distribution over the $2^{50,000}$ possible outcomes, then we'd end up with a $(2^{50,000} - 1)$ dimensional parameter vector.

↳ This is clearly too many parameters.

⇒ To model $P(x|y)$, we will therefore make a very strong assumption.

↳ We will assume that the x_i 's are conditionally independent given y .

⇒ This assumption is called the Naive Bayes (NB) Assumption.

⇒ The resulting algorithm is called the Naive Bayes Classifier.

$$\begin{aligned}
 & P(x_1, \dots, x_{50,000} | y) \\
 &= P(x_1 | y) P(x_2 | y, x_1) P(x_3 | y, x_1, x_2) \\
 &\quad \dots P(x_{50,000} | y, x_1, x_2, \dots, x_{49,999}) \\
 &= P(x_1 | y) P(x_2 | y) \dots P(x_{50,000} | y) \\
 &= \prod_{i=1}^n P(x_i | y)
 \end{aligned}$$

⇒ Even though the Naive Bayes assumption is an extremely strong assumption, the resulting algorithm works well on many problems.

⇒ Our model is parameterized by $\phi_{i|y=1}$, $\phi_{i|y=0}$ & ϕ_y

$\phi_{i|y=1}$
 \downarrow
 $P(x_i=1 | y=1)$

$\phi_{i|y=0}$
 \downarrow
 $P(x_i=1 | y=0)$

ϕ_y
 \downarrow
 $P(y=1)$

①
⇒ Given the training set $\{(x^{(i)}, y^{(i)}) ; i=1, \dots, m\}$
We can write joint likelihood of the data:

$$\mathcal{L}(\phi_y, \phi_{j|y=0}, \phi_{j|y=1}) = \prod_{i=1}^m P(x^{(i)}, y^{(i)})$$

⇒ Maximizing this with respect to $\phi_y, \phi_{j|y=0}$ & $\phi_{j|y=1}$ gives the maximum likelihood estimates:

$$\phi_{j|y=1} = \frac{\sum_{i=1}^m 1\{x_j^{(i)}=1 \wedge y^{(i)}=1\}}{\sum_{i=1}^m 1\{y^{(i)}=1\}}$$

$$\phi_{j|y=0} = \frac{\sum_{i=1}^m 1\{x_j^{(i)}=1 \wedge y^{(i)}=0\}}{\sum_{i=1}^m 1\{y^{(i)}=0\}}$$

$$\phi_y = \frac{\sum_{i=1}^m 1\{y^{(i)}=1\}}{m}$$

⇒ To make a prediction on a new example with features x , we then simply calculate:

$$P(y=1|x) = \frac{P(x|y=1) P(y=1)}{P(x)}$$

$$= \frac{\left(\prod_{i=1}^n P(x_i|y=1) \right) P(y=1)}{\left(\prod_{i=1}^n P(x_i|y=1) \right) P(y=1) + \left(\prod_{i=1}^n P(x_i|y=0) \right) P(y=0)}$$

$$+ \left(\prod_{i=1}^n P(x_i|y=0) \right) P(y=0)$$

⇒ and pick whichever class have the higher Posterior probability.

⇒ Generalization of Naive Bayes where X_i can take values in $\{1, 2, \dots, K_i\}$ is straight forward.

↳ Here we simply model $P(X_i | y)$ as multinomial rather than Bernoulli:

⇒ When the original, continuous-valued attributes are not well modeled by a multivariate normal distribution, discretizing the feature and using Naive Bayes (instead of GDA) will often result in a better classifier.

★ Laplace smoothing

⇒ Statistically ^{its} a bad idea to estimate the probability of some event to be zero just because you haven't seen it before in your finite training set.

⇒ Take the problem of estimating the mean of multinomial random variable Z taking values in $\{1, \dots, K\}$.

⇒ We can parameterize our multinomial with $\phi_i = P(Z=i)$

⇒ Given a set of m independent observations $\{Z^{(1)}, \dots, Z^{(m)}\}$, the maximum likelihood estimate are given by

$$\phi_i = \frac{\sum_{j=1}^m 1\{Z^{(j)}=i\}}{m}$$

⇒ If we use maximum likelihood estimates, then some of the ϕ_j 's might end up zero.

⇒ To avoid this we can use Laplace smoothing which replaces the above estimate with

$$\phi_j = \frac{\sum_{i=1}^m 1\{z^{(i)} = j\} + 1}{m + K}$$

⇒ Returning back to our Naive Bayes classifier, with Laplace smoothing we therefore obtain the following estimate of the parameters:

$$\phi_{j|y=1} = \frac{\sum_{i=1}^m 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 1\} + 1}{\sum_{i=1}^m 1\{y^{(i)} = 1\} + 2}$$

$$\phi_{j|y=0} = \frac{\sum_{i=1}^m 1\{x_j^{(i)} = 1 \wedge y^{(i)} = 0\} + 1}{\sum_{i=1}^m 1\{y^{(i)} = 0\} + 2}$$

★ Event models for text classification

⇒ Naive Bayes works well for many classification problems, for text classification there is a related model that does even better.

⇒ Naive Bayes as presented uses

{ Multi-variate Bernoulli event model }

⇒ Here's a different model, called the multinomial event model

⇒ We will use a different notation and set of features for representing emails.

⇒ Let x_i denotes the identity of the i^{th} word in the email.

$$x_i \in \{1, \dots, |V|\}$$

where $|V| \rightarrow$ size of vocabulary

⇒ An email of n words is now represented by a vector (x_1, x_2, \dots, x_n) of length n .

↳ n can vary for different documents.

⇒ On the multinomial event model, we assume that the way an email is generated is via a random

Process:

① → Spam/non-spam is first determined

② → Then, the sender of the email writes the email by first generating x_1 from some multinomial distribution over words $P(x_1 | y)$

③ → Next the second word is x_2 is chosen independently of x_1 , but from the same multinomial distribution, and similarly for x_3, x_4 , and so on.

③ \hookrightarrow Thus the overall probability of a message is given by $P(y) \prod_{i=1}^n P(x_i | y)$

\Rightarrow The parameters of our new model are:

$$\phi_y = P(y)$$

$$\phi_{i|y=1} = P(x_i = i | y = 1)$$

$$\phi_{i|y=0} = P(x_i = i | y = 0)$$

$\left. \begin{array}{l} \text{for any } i \end{array} \right\}$

\Rightarrow If we are given a training set

$$\{(x^{(i)}, y^{(i)}); i = 1, \dots, m\}$$

$$\text{where } x^{(i)} = (x_1^{(i)}, x_2^{(i)}, \dots, x_{n_i}^{(i)})$$

\Rightarrow The likelihood of the data is given by

$$\mathcal{L}(\phi, \phi_{i|y=0}, \phi_{i|y=1}) = \prod_{i=1}^m P(x^{(i)}, y^{(i)})$$

$$= \prod_{i=1}^m \left(\prod_{j=1}^{n_i} P(x_j^{(i)} | y; \phi_{i|y=0}, \phi_{i|y=1}) \right) P(y^{(i)}; \phi_y)$$

\Rightarrow Maximizing this yields the maximum likelihood estimates of the parameters:

⇒ Maximizing this yields the maximum likelihood estimates of the parameters:

$$\phi_{k/y=1} = \frac{\sum_{i=1}^m \sum_{j=1}^{n_i} 1\{x_j^{(i)} = k \wedge y^{(i)} = 1\} + 1}{\sum_{i=1}^m 1\{y^{(i)} = 1\} n_i + |V|}$$

$$\phi_{k/y=0} = \frac{\sum_{i=1}^m \sum_{j=1}^{n_i} 1\{x_j^{(i)} = k \wedge y^{(i)} = 0\} + 1}{\sum_{i=1}^m 1\{y^{(i)} = 0\} n_i + |V|}$$

$$\phi_y = \frac{\sum_{i=1}^m 1\{y^{(i)} = 1\}}{m}$$

} Laplace Smoothing
Added

⇒ While not necessarily the very best classification algorithm, the Naive Bayes Classifier often works surprisingly well.

↳ It is often also a very good "first thing to try" given its simplicity & ease of implementation.