(11) Detection & Segmentation & Localization

Semantic Segmentation | Classification + Localization | Object Detection | Instance Segmentation

★ Semantic Segmentation

- Image (Input) ⟶ decision of category for every Pixel in the image (Output)

⇒ Don't differentiate instances, only care about pixels

Idea1: Sliding Window

↳ Problem: Very inefficient! Not reusing shared features between overlapping patches.

Idea 2: Fully Convolutional network

⇒ Design a network as a bunch of convolutional layers to make predictions for Pixels all at once.

⇒ Design network as a bunch of convolutional layers, with downsampling & upsampling inside the network.

In-Network upsampling: Unpooling

1	2
3	4

Nearest Neighbor

1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

1	2
3	4

Bed of Nails

1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

Max Pooling

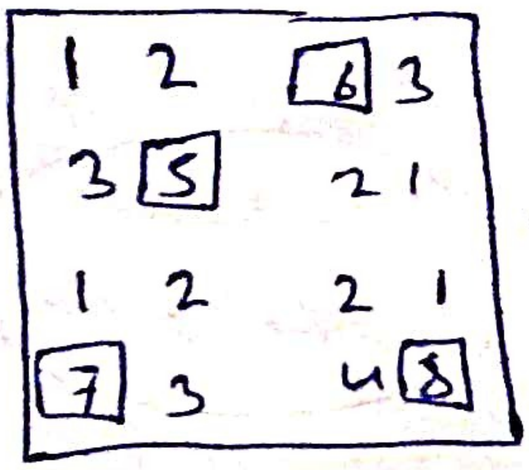

Remember which element was max!

Max Unpooling

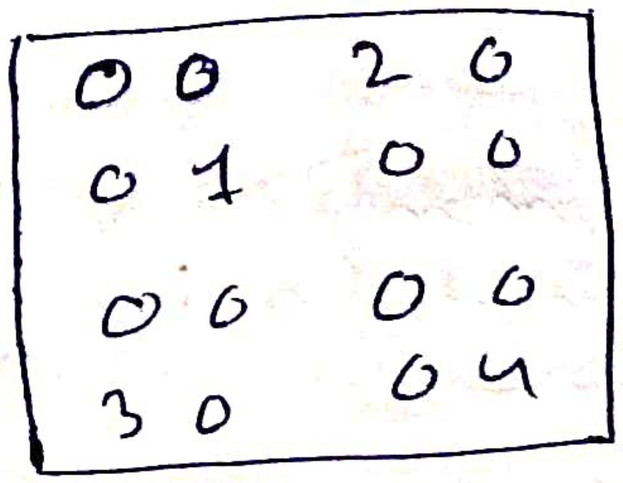

Use position from Pooling layer.

Learnable Upsampling: Transpose Convolution

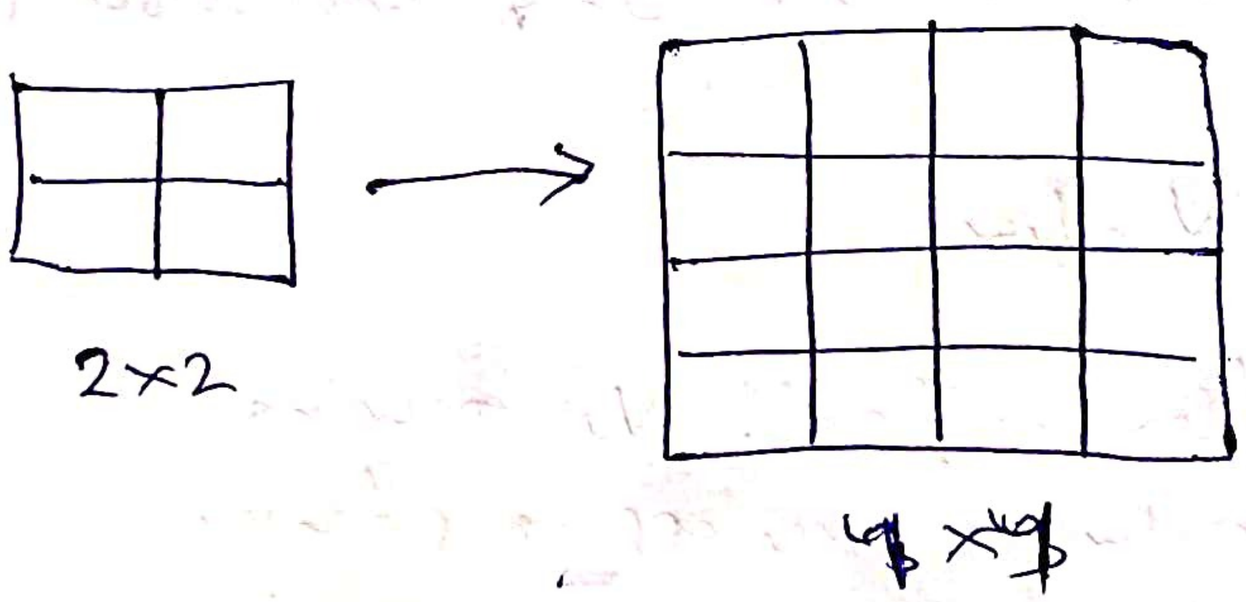

⇒ It uses filter for upsampling.

★ Classification + Localization

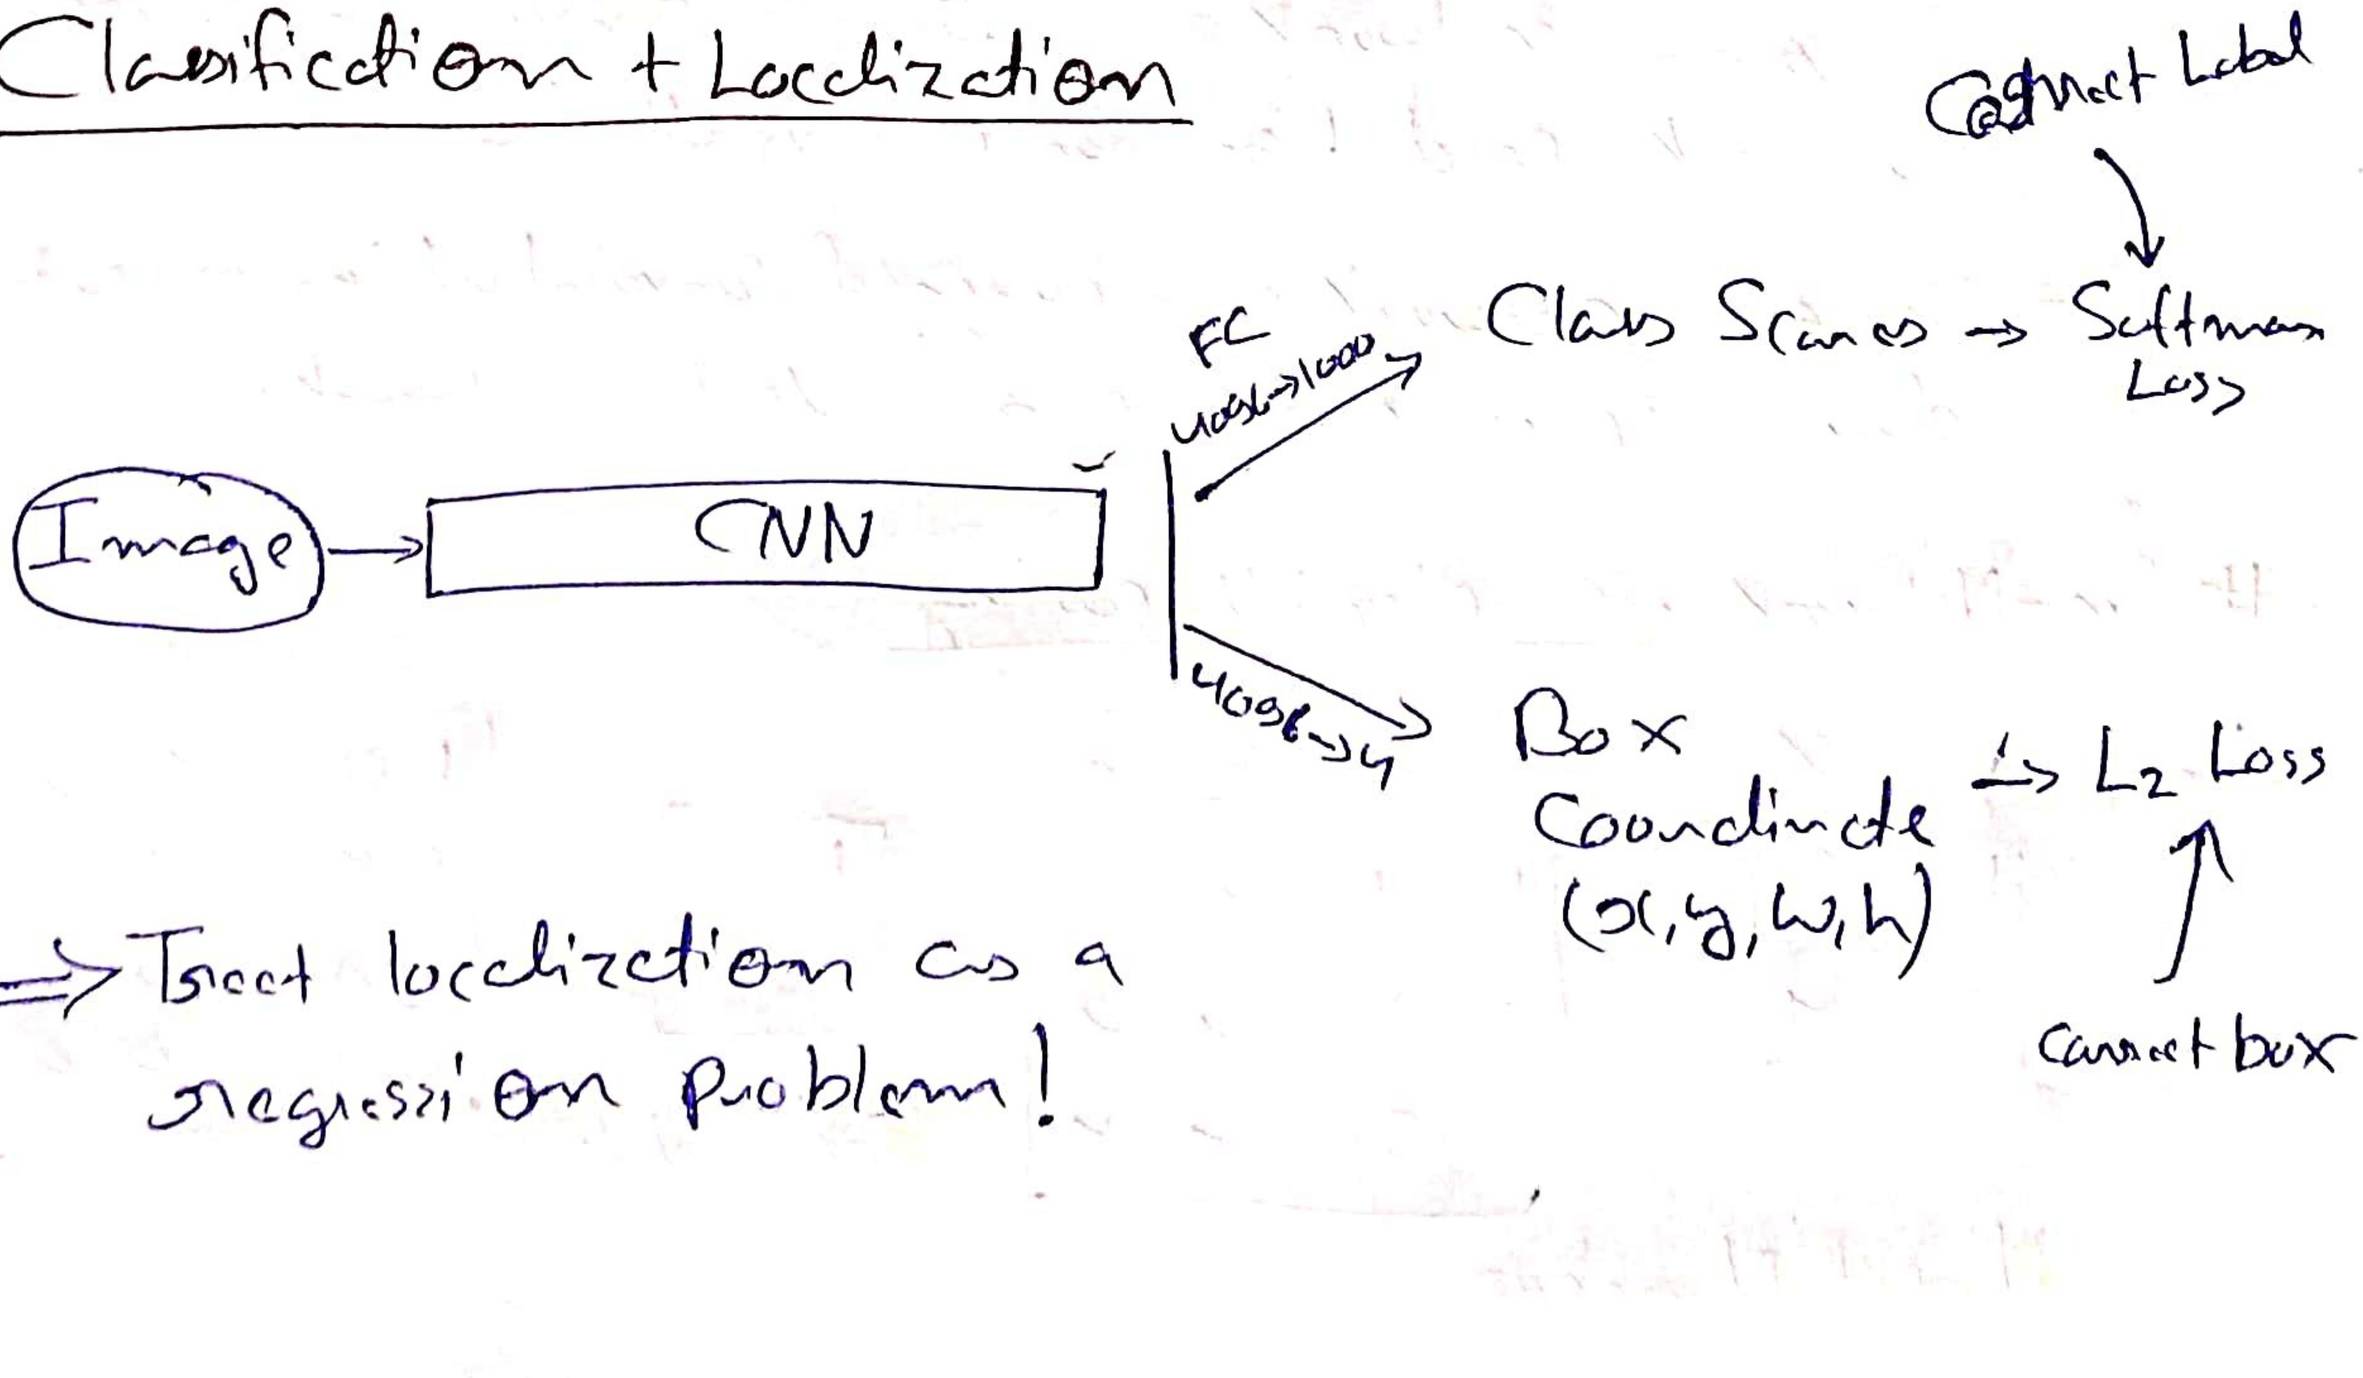

⇒ Treat localization as a regression problem!

* Object Detection

⇒ Given an image draw bounding box around all the class in the image and give the class label.

⇒ Each image needs a different number of output.

⇒ One approach: Sliding Window

↳ Problem: Need to apply CNN to huge number of locations & scales, very computationally expensive!

≠ Region Proposal

⇒ Find "blobby" image regions that are likely to contain Object.

R-CNN

{Region Based Convolutional Neural Networks}

Fast R-CNN

Faster R-CNN

↳ Jointly train with 4 losses:

① → RPN classify object / not object
② → RPN regress box coordinate
③ → Find classification score (object class)
④ → Find box coordinate

⊜ Detection without Proposals: YOLO / SSD

{You Only Look once} ←

{Single-Shot Detection} ←

⇒ Within each grid cell:

→ Regress from each of the B base boxes to a final box with 5 numbers.
(dx dy dh dw, confidence)

→ Predict score for each of C classes.
(Including background as a class)

★ Instance Segmentation

(Mask R-CNN)